

**Київський національний університет імені Тараса Шевченка
радіофізичний факультет**

Лабораторна робота № 3

Тема:

«Дослідження оптимізації коду з використанням векторних
розширень CPU»

Роботу виконав
студент 3 курсу
Комп'ютерна Інженерія
Стецюк Владислав

Київ 2019

Хід роботи

1. Отримайте доступ на обчислювальний кластер для роботи з Intel Compiler

Certificate details

Distinguished name (DN):	/C=UA/O=KNU/OU=People/L=FRECS/CN=Vladislav Stetsiuk
Serial number:	DAC746B1B36BCFAC
Valid from:	Wed, 10 Apr 2019 09:02:00
Valid to:	Mon, 07 Oct 2019 09:02:00
Certificate type:	Personal User Certificate
Certificate hash:	e109bb6a

2. Завантажте файли Intel® C++ Compiler - Using Auto-Vectorization Tutorial (<https://software.intel.com/en-us/product-code-samples?topic=20813>) на свій комп'ютер та в домашню директорію користувача обчислювального кластеру.

```
KNU: :s3 [tb218 ~]$ wget https://software.intel.com/sites/default/files/vec_samples_C_lin_20170911.tgz
```

```
[tb218@plus7 ~]$ tar -zxvf vec_samples_C_lin_20170911.tgz
./vec_samples/
./vec_samples/license.txt
./vec_samples/src/
./vec_samples/src/Multiply.h
./vec_samples/src/Driver.c
./vec_samples/src/Multiply.c
./vec_samples/build.bat
./vec_samples/vec_samples_2017.sln
./vec_samples/Makefile
```

```
[tb218@plus7 ~]$ cd vec_samples
[tb218@plus7 vec_samples]$ ls
build.bat      msvs2017      vec_samples_2013.sln
license.txt    readme.html   vec_samples_2015.sln
Makefile       resources     vec_samples_2017.sln
msvs2013      src
msvs2015      tutorial
```

3. Використовуючи інструкції в readme.html ознайомтесь та виконайте Tutorial на обчислювальному кластері.

```
KNU: :s3 [tb218 vec_samples]$ ls
Makefile      msvs2017      vec_samples_2013.sln
build.bat     readme.html   vec_samples_2015.sln
license.txt   resources     vec_samples_2017.sln
msvs2013      src
msvs2015      tutorial
KNU: :s3 [tb218 vec_samples]$ icc -O1 -std=c99 src/Multiply.c src
/Driver.c -o MatVector
KNU: :s3 [tb218 vec_samples]$ ./MatVector

ROW:101 COL: 101
Execution time is 12.142 seconds
GigaFlops = 1.680282
Sum of result = 195853.999899
KNU: :s3 [tb218 vec_samples]$ icc -std=c99 -O2 -D NOFUNCCALL -qop
t-report=1 -qopt-report-phase=vec src/Multiply.c src/Driver.c -o M
atVector
icc: remark #10397: optimization reports are generated in *.optrpt
files in the output location
KNU: :s3 [tb218 vec_samples]$ ./MatVector

ROW:101 COL: 101
Execution time is 4.123 seconds
GigaFlops = 4.948244
Sum of result = 195853.999899
KNU: :s3 [tb218 vec_samples]$ cat Multiply.optrpt
Intel(R) Advisor can now assist with vectorization and show optimi
zation
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for detail
s.
```

Begin optimization report for: matvec(int, int, double (*)[*], dou
ble *, double *)

Report from: Vector optimizations [vec]

```
LOOP BEGIN at src/Multiply.c(37,5)
remark #25460: No loop optimizations reported

LOOP BEGIN at src/Multiply.c(49,9)
remark #25460: No loop optimizations reported
LOOP END

LOOP BEGIN at src/Multiply.c(49,9)
```

```
GNU: :s3 [tb218 vec_samples]$ icc -std=c99 -qopt-report=2 -qopt-report-phase=vec -D NOALIAS src/Multiply.c src/Driver.c -o MatVector
icc: remark #10397: optimization reports are generated in *.optrpt files in the output location
GNU: :s3 [tb218 vec_samples]$ ./MatVector
```

```
ROW:101 COL: 101
Execution time is 8.315 seconds
GigaFlops = 2.453751
Sum of result = 195853.999899
```

```
GNU: :s3 [tb218 vec_samples]$ cat Multiply.optrpt
Intel(R) Advisor can now assist with vectorization and show optimization
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for details.
```

```
Begin optimization report for: matvec(int, int, double (*)[*], double *__restrict__, double *)
```

```
Report from: Vector optimizations [vec]
```

```
LOOP BEGIN at src/Multiply.c(37,5)
remark #15542: loop was not vectorized: inner loop was already vectorized
```

```
LOOP BEGIN at src/Multiply.c(49,9)
<Peeled loop for vectorization>
LOOP END
```

```
LOOP BEGIN at src/Multiply.c(49,9)
remark #15300: LOOP WAS VECTORIZED
LOOP END
```

```
LOOP BEGIN at src/Multiply.c(49,9)
<Alternate Alignment Vectorized Loop>
LOOP END
```

```
LOOP BEGIN at src/Multiply.c(49,9)
<Remainder loop for vectorization>
LOOP END
```

```
LOOP END
```

GNU: :s3 [tb218 vec_samples]\$ ==> PBS: job killed: walltime 1834 exceeded limit 1800

GNU: :s3 [tb218 vec_samples]\$ icc -std=c99 -qopt-report=4 -qopt-report-phase=vec -D NOALIAS -D ALIGNED src/Multiply.c src/Driver.c -o MatVector

icc: remark #10397: optimization reports are generated in *.optrpt files in the output location

GNU: :s3 [tb218 vec_samples]\$ cat Multiply.optrpt

Intel(R) Advisor can now assist with vectorization and show optimization

report messages with your source code.

See "<https://software.intel.com/en-us/intel-advisor-xe>" for details.

Intel(R) C Intel(R) 64 Compiler for applications running on Intel(R) 64, Version 18.0.5.274 Build 20180823

Compiler options: -std=c99 -qopt-report=4 -qopt-report-phase=vec -D NOALIAS -D ALIGNED -o MatVector

Begin optimization report for: matvec(int, int, double (*)[*], double *__restrict__, double *)

Report from: Vector optimizations [vec]

LOOP BEGIN at src/Multiply.c(37,5)

remark #15542: loop was not vectorized: inner loop was already vectorized

LOOP BEGIN at src/Multiply.c(49,9)

remark #15388: vectorization support: reference a[i][j] has aligned access [src/Multiply.c(50,21)]

```

    remark #15305: vectorization support: vector length 2
    remark #15399: vectorization support: unroll factor set to 4
    remark #15309: vectorization support: normalized vectorizati
on overhead 0.594
    remark #15300: LOOP WAS VECTORIZED
    remark #15448: unmasked aligned unit stride loads: 2
    remark #15475: --- begin vector cost summary ---
    remark #15476: scalar cost: 10
    remark #15477: vector cost: 4.000
    remark #15478: estimated potential speedup: 2.410
    remark #15488: --- end vector cost summary ---
LOOP END

```

```

LOOP BEGIN at src/Multiply.c(49,9)
<Remainder loop for vectorization>
    remark #15388: vectorization support: reference a[i][j] has
aligned access    [ src/Multiply.c(50,21) ]
    remark #15388: vectorization support: reference x[j] has ali
gned access    [ src/Multiply.c(50,31) ]
    remark #15335: remainder loop was not vectorized: vectorizat
ion possible but seems inefficient. Use vector always directive or
-vec-threshold0 to override
    remark #15305: vectorization support: vector length 2
    remark #15309: vectorization support: normalized vectorizati
on overhead 2.417
    LOOP END
LOOP END

```

```

=====
=====

```

```

KNU:  :s1 [tb218 vec_samples]$ icc -std=c99 -qopt-report=2 -qopt-r
eport-phase=vec -D NOALIAS -D ALIGNED -ipo src/Multiply.c src/Driv
er.c -o MatVector

```

```
GNU: :s1 [tb218 vec_samples]$ cat ipo_out.optrpt
Intel(R) Advisor can now assist with vectorization and show optimization
report messages with your source code.
See "https://software.intel.com/en-us/intel-advisor-xe" for details.
```

```
Begin optimization report for: main()
```

```
Report from: Vector optimizations [vec]
```

```
LOOP BEGIN at src/Driver.c(152,16)
```

```
remark #15542: loop was not vectorized: inner loop was already
vectorized
```

```
LOOP BEGIN at src/Multiply.c(37,5) inlined into src/Driver.c(150,9)
```

```
remark #15542: loop was not vectorized: inner loop was already
vectorized
```

```
LOOP BEGIN at src/Multiply.c(49,9) inlined into src/Driver.c
(150,9)
```

```
remark #15300: LOOP WAS VECTORIZED
```

```
LOOP END
```

```
LOOP BEGIN at src/Multiply.c(49,9) inlined into src/Driver.c
(150,9)
```

```
<Remainder loop for vectorization>
```

```
remark #15335: remainder loop was not vectorized: vectorization
possible but seems inefficient. Use vector always directive
```

```

LOOP BEGIN at src/Driver.c(74,5) inlined into src/Driver.c(159,5)
<Remainder loop for vectorization>
LOOP END
=====
=====

Begin optimization report for: init_matrix(int, int, double, double (*)[102])

    Report from: Vector optimizations [vec]

LOOP BEGIN at src/Driver.c(47,5)
    remark #15542: loop was not vectorized: inner loop was already
    vectorized

    LOOP BEGIN at src/Driver.c(48,9)
        remark #15300: LOOP WAS VECTORIZED
    LOOP END

    LOOP BEGIN at src/Driver.c(48,9)
        <Remainder loop for vectorization>
    LOOP END
LOOP END

LOOP BEGIN at src/Driver.c(53,9)
    remark #15300: LOOP WAS VECTORIZED
LOOP END

LOOP BEGIN at src/Driver.c(53,9)
<Remainder loop for vectorization>
LOOP END
=====
=====

KNU:  :s1 [tb218 vec_samples]$ ./MatVector

ROW:101 COL: 102
Execution time is 3.995 seconds
GigaFlops = 5.107272
Sum of result = 195853.999899

```

4. Оберіть будь-яку неінтерактивну консольну програму мовою C/C++ (унікальну в межах групи, в гуглі більше ніж 50 програм)

a) Напишіть сценарій, що:

- i) Компілює програму з різними оптимізаціями (-O) та виміряйте час її роботи. Якщо час досить малий - вимірюйте час роботи 1000 (чи 1000000) запусків алгоритму в циклі. Час роботи можна виміряти утилітою time.
- ii) Отримує перелік всіх розширень процесору що підтримуються
- iii) Для кожного розширення компілює Intel-компілятором окремий варіант оптимізованого коду (наприклад -x SSE2)
- iv) Вимірює час виконання кожного варіанта оптимізованої програми

```
GNU: :s1 [tb218 ~]$ cat date.c
#include <stdio.h>
int main()
{
    int daysInCurrentFebruary = 29;
    int daysInJanuary = 31;
    int daysInFebruary = daysInCurrentFebruary;
    int daysInMarch = 31;
    int daysInApril = 30;
    int daysInMay = 31;
    int daysInJune = 30;
    int daysInJuly = 31;
    int daysInAugust = 31;
    int daysInSeptember = 30;
    int daysInOctober = 31;
    int daysInNovember = 30;
    int daysInDecember = 31;
    int daysInFirstHalf = daysInJanuary + daysInFebruary + daysInMarch + daysInApril + daysInMay + daysInJune;
    int daysInSecondHalf = daysInJuly + daysInAugust + daysInSeptember + daysInOctober + daysInNovember + daysInDecember;
    printf("Days in the first half of the current year: %d\n", daysInFirstHalf);
    printf("Days in the second half of the current year: %d\n", daysInSecondHalf);
    printf("Days in the current year: %d\n", daysInFirstHalf + daysInSecondHalf);
    return 0;
}
GNU: :s1 [tb218 ~]$

GNU: :s1 [tb218 ~]$ gcc date.c
GNU: :s1 [tb218 ~]$ ls
ITAC  advisor  compiler_c  date.c  inspector  ipsxe2019_samples_lin_20190327.tgz  mkl  tbb
a.out cluster_checker compiler_f  index.html  ipp  licensing  pstl  vtune_amplifier
GNU: :s1 [tb218 ~]$ ./a.out
Days in the first half of the current year: 182
Days in the second half of the current year: 184
Days in the current year: 366
```

```
GNU: :s1 [tb218 ~]$ time gcc date.c
```

```
real    0m0.075s
user    0m0.019s
sys     0m0.018s
```

```
GNU: :s3 [tb218 ~]$ cat 1.sh
#!/bin/bash
```

```
flags=( "sse4.2" "sse4.1" "sse3" "sse2" "ssse3" "avx" )
```

```
for i in "${flags[@]}; do
    for j in {1..3}; do
        gcc -O$j -m$i date.c -o Temp
        echo
        echo $i " " $j
        time `for i in {0..1000}; do ./Temp; done`
    done
done
```

—

```
GNU: :s1 [tb218 ~]$ ./1.sh
```

```
sse4.2 1
```

```
./1.sh: line 10: Days: command not found
```

```
real    0m1.321s
```

```
user    0m0.200s
```

```
sys     0m1.111s
```

```
sse4.2 2
```

```
./1.sh: line 10: Days: command not found
```

```
real    0m1.400s
```

```
user    0m0.193s
```

```
sys     0m1.199s
```

```
sse4.2 3
```

```
./1.sh: line 10: Days: command not found
```

```
real    0m1.403s
```

```
user    0m0.156s
```

```
sys     0m1.237s
```

```
sse4.1 1
```

```
./1.sh: line 10: Days: command not found
```

```
real    0m1.394s
```

```
user    0m0.178s
```

```
sys     0m1.208s
```

```

ssse3  3
./1.sh: line 10: Days: command not found

real    0m0.908s
user    0m0.313s
sys     0m0.588s

avx    1
./1.sh: line 10: Days: command not found

real    0m0.911s
user    0m0.297s
sys     0m0.606s

avx    2
./1.sh: line 10: Days: command not found


real    0m0.911s
user    0m0.309s
sys     0m0.595s

avx    3
./1.sh: line 10: Days: command not found

real    0m0.917s
user    0m0.300s
sys     0m0.609s

```

	sse 4.2	sse 4.1	sse 3	sse 2	ssse 3	avx
O1	1,321	1,394	0,938	0,909	0,910	0,911
O2	1,400	1,406	0,911	0,909	0,911	0,911
O3	1,403	1,090	0,912	0,909	0,908	0,917
Average	1,3746	1,2966	0,920	0,909	0,9096	0,913



- b. Запустіть задачу **в планувальник** обчислювального кластеру 5 разів (для статистики на різних нодах)

```
[tb218@plus7 ~]$ qsub -N MyJob -l nodes=1:ppn=1,walltime=00:30:00
1.sh
2789321
[tb218@plus7 ~]$ qsub -N MyJob -l nodes=1:ppn=1,walltime=00:30:00
1.sh
2789322
```

Як бачимо з отриманих даних, найшвидше програма виконується у випадку компіляції з третім методом оптимізації та розширенням процесора sse2.

Висновок: У процесі виконання лабораторної роботи було проведено ознайомлення з обчислювальним кластером та методами оптимізації виконання програми на C/C++ . За результатами виконання було написано сценарії та звіт з відповідними скріншотами.